

m8w4.rmd

raymond

June 1, 2017

Prediction Assignment Writeup

This project asks to analyze data by personal activity devices. We first load the data

```
training <- read.csv(url("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"))
testing <- read.csv(url("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"))
```

Now, we load the necessary packages, set the seed (for it to be reproducible) and a look at the data

```
## 'data.frame':    19622 obs. of  160 variables:
## $ X                      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ user_name              : Factor w/ 6 levels "adelmo","carlitos",...: 2 2 2 2 2 2 2 2 2 2 ...
## $ raw_timestamp_part_1   : int  1323084231 1323084231 1323084231 1323084232 1323084232 1323084232 1323084232 1323084232 1323084232 ...
## $ raw_timestamp_part_2   : int  788290 808298 820366 120339 196328 304277 368296 440390 484323 484434 ...
## $ cvtd_timestamp        : Factor w/ 20 levels "02/12/2011 13:32",...: 9 9 9 9 9 9 9 9 9 9 ...
## $ new_window             : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ num_window             : int  11 11 11 12 12 12 12 12 12 12 ...
## $ roll_belt              : num  1.41 1.41 1.42 1.48 1.48 1.45 1.42 1.42 1.43 1.45 ...
## $ pitch_belt             : num  8.07 8.07 8.07 8.05 8.07 8.06 8.09 8.13 8.16 8.17 ...
## $ yaw_belt               : num  -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -9 ...
## $ total_accel_belt       : int  3 3 3 3 3 3 3 3 3 3 ...
## $ kurtosis_roll_belt     : Factor w/ 397 levels "", "-0.016850",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ kurtosis_pitch_belt   : Factor w/ 317 levels "", "-0.021887",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ kurtosis_yaw_belt     : Factor w/ 2 levels "", "#DIV/0!": 1 1 1 1 1 1 1 1 1 1 ...
## $ skewness_roll_belt    : Factor w/ 395 levels "", "-0.003095",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ skewness_roll_belt.1  : Factor w/ 338 levels "", "-0.005928",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ skewness_yaw_belt     : Factor w/ 2 levels "", "#DIV/0!": 1 1 1 1 1 1 1 1 1 1 ...
## $ max_roll_belt         : num  NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_belt        : int  NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_belt          : Factor w/ 68 levels "", "-0.1", "-0.2",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ min_roll_belt         : num  NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_belt        : int  NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_belt          : Factor w/ 68 levels "", "-0.1", "-0.2",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ amplitude_roll_belt   : num  NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_pitch_belt  : int  NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_belt    : Factor w/ 4 levels "", "#DIV/0!", "0.00",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ var_total_accel_belt  : num  NA NA NA NA NA NA NA NA NA NA ...
## $ avg_roll_belt         : num  NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_belt      : num  NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_belt         : num  NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_belt        : num  NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_belt     : num  NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_belt        : num  NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_belt          : num  NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_belt       : num  NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_belt          : num  NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_belt_x           : num  0 0.02 0 0.02 0.02 0.02 0.02 0.02 0.02 0.03 ...
## $ gyros_belt_y           : num  0 0 0 0 0.02 0 0 0 0 0 ...
## $ gyros_belt_z           : num  -0.02 -0.02 -0.02 -0.03 -0.02 -0.02 -0.02 -0.02 -0.02 0 ...
## $ accel_belt_x          : int  -21 -22 -20 -22 -21 -21 -22 -22 -20 -21 ...
## $ accel_belt_y          : int  4 4 5 3 2 4 3 4 2 4 ...
## $ accel_belt_z          : int  22 22 23 21 24 21 21 21 24 22 ...
## $ magnet_belt_x         : int  -3 -7 -2 -6 -6 0 -4 -2 1 -3 ...
```

```

## $ magnet_belt_y      : int  599 608 600 604 600 603 599 603 602 609 ...
## $ magnet_belt_z      : int  -313 -311 -305 -310 -302 -312 -311 -313 -312 -308 ...
## $ roll_arm           : num  -128 -128 -128 -128 -128 -128 -128 -128 -128 -128 ...
## $ pitch_arm          : num   22.5 22.5 22.5 22.1 22.1 22 21.9 21.8 21.7 21.6 ...
## $ yaw_arm            : num  -161 -161 -161 -161 -161 -161 -161 -161 -161 -161 ...
## $ total_accel_arm    : int   34 34 34 34 34 34 34 34 34 34 ...
## $ var_accel_arm      : num   NA NA NA NA NA NA NA NA NA NA ...
## $ avg_roll_arm       : num   NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_arm    : num   NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_arm       : num   NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_arm      : num   NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_arm   : num   NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_arm      : num   NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_arm        : num   NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_arm     : num   NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_arm        : num   NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_arm_x        : num    0 0.02 0.02 0.02 0 0.02 0 0.02 0.02 0.02 ...
## $ gyros_arm_y        : num    0 -0.02 -0.02 -0.03 -0.03 -0.03 -0.03 -0.02 -0.03 -0.03
...
## $ gyros_arm_z        : num   -0.02 -0.02 -0.02 0.02 0 0 0 0 -0.02 -0.02 ...
## $ accel_arm_x        : int  -288 -290 -289 -289 -289 -289 -289 -289 -288 -288 ...
## $ accel_arm_y        : int   109 110 110 111 111 111 111 111 109 110 ...
## $ accel_arm_z        : int  -123 -125 -126 -123 -123 -122 -125 -124 -122 -124 ...
## $ magnet_arm_x       : int  -368 -369 -368 -372 -374 -369 -373 -372 -369 -376 ...
## $ magnet_arm_y       : int   337 337 344 344 337 342 336 338 341 334 ...
## $ magnet_arm_z       : int   516 513 513 512 506 513 509 510 518 516 ...
## $ kurtosis_roll_arm  : Factor w/ 330 levels "", "-0.02438",...: 1 1 1 1 1 1 1 1 1 1
...
## $ kurtosis_pitch_arm : Factor w/ 328 levels "", "-0.00484",...: 1 1 1 1 1 1 1 1 1 1
...
## $ kurtosis_yaw_arm   : Factor w/ 395 levels "", "-0.01548",...: 1 1 1 1 1 1 1 1 1 1
...
## $ skewness_roll_arm  : Factor w/ 331 levels "", "-0.00051",...: 1 1 1 1 1 1 1 1 1 1
...
## $ skewness_pitch_arm : Factor w/ 328 levels "", "-0.00184",...: 1 1 1 1 1 1 1 1 1 1
...
## $ skewness_yaw_arm   : Factor w/ 395 levels "", "-0.00311",...: 1 1 1 1 1 1 1 1 1 1
...
## $ max_roll_arm       : num   NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_arm      : num   NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_arm        : int    NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_arm       : num   NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_arm      : num   NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_arm        : int    NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_arm : num   NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_pitch_arm : num   NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_arm  : int    NA NA NA NA NA NA NA NA NA NA ...
## $ roll_dumbbell      : num   13.1 13.1 12.9 13.4 13.4 ...
## $ pitch_dumbbell     : num  -70.5 -70.6 -70.3 -70.4 -70.4 ...
## $ yaw_dumbbell       : num  -84.9 -84.7 -85.1 -84.9 -84.9 ...
## $ kurtosis_roll_dumbbell : Factor w/ 398 levels "", "-0.0035", "-0.0073",...: 1 1 1 1 1 1 1 1
1 1 1 ...
## $ kurtosis_pitch_dumbbell : Factor w/ 401 levels "", "-0.0163", "-0.0233",...: 1 1 1 1 1 1 1 1
1 1 1 ...
## $ kurtosis_yaw_dumbbell : Factor w/ 2 levels "", "#DIV/0!": 1 1 1 1 1 1 1 1 1 1 ...
## $ skewness_roll_dumbbell : Factor w/ 401 levels "", "-0.0082", "-0.0096",...: 1 1 1 1 1 1 1 1
1 1 1 ...
## $ skewness_pitch_dumbbell : Factor w/ 402 levels "", "-0.0053", "-0.0084",...: 1 1 1 1 1 1 1 1

```

```

1 1 1 ...
## $ skewness_yaw_dumbbell : Factor w/ 2 levels "", "#DIV/0!": 1 1 1 1 1 1 1 1 1 ...
## $ max_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_dumbbell : Factor w/ 73 levels "", "-0.1", "-0.2",...: 1 1 1 1 1 1 1 1 1 1
...
## $ min_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_dumbbell : Factor w/ 73 levels "", "-0.1", "-0.2",...: 1 1 1 1 1 1 1 1 1 1
...
## $ amplitude_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA ...
## [list output truncated]

```

Cleaning the data

Now, we commence the cleaning of data, which includes deleting off unrelated and useless columns (containing a lot of NAs). First off is the training dataset.

```

cols_na <- nearZeroVar(training) #cols with little/no variance
training <- training[, -cols_na]

keep_index <- !sapply(training, function(x) any(is.na(x))) #del cols containing NAs
training <- training[, keep_index]
keep_index <- sapply(colnames(training), function(x) !grepl("X|time|window",x))
# ^ remove cols with labeling functions
training <- training[, keep_index]
dim(training)

```

```
## [1] 19622 54
```

Now, we do the same filtering to the testing dataset.

```

keep_index <- !sapply(testing, function(x) any(is.na(x)))
testing <- testing[, keep_index]
keep_index <- sapply(colnames(testing), function(x) !grepl("X|time|window",x))
testing <- testing[, keep_index];
#remove problem_id col
idx1 <- which(colnames(testing)=="problem_id")
testing <- testing[,-idx1]
dim(testing)

```

```
## [1] 20 53
```

Machine Learning

Now, we splice the training dataset so we have a 'train' and 'test' data from the training dataset

```

index_train <- createDataPartition(training$classe, p = 0.7, list=FALSE)
training1 <- training[index_train, ]
testing1 <- training[-index_train, ]

```

Side note:-

```
control <- trainControl(method = "cv", number = 5)
```

We set the number of cross validation to 5 instead of the default 10 to save computation time. Also my computer ran out of memory with the default 10.

LDA

First, we try Linear Discriminant Analysis (LDA)

```
modFit_lda <- train(classe ~., data=training1, method="lda")  
print(modFit_lda, digits = 4)
```

```
## Linear Discriminant Analysis  
##  
## 13737 samples  
##    53 predictor  
##    5 classes: 'A', 'B', 'C', 'D', 'E'  
##  
## No pre-processing  
## Resampling: Bootstrapped (25 reps)  
## Summary of sample sizes: 13737, 13737, 13737, 13737, 13737, 13737, ...  
## Resampling results:  
##  
##   Accuracy   Kappa  
##   0.7355    0.6648
```

```
predict_lda <- predict(modFit_lda, testing1)  
(conf_lda <- confusionMatrix(testing1$classe, predict_lda))
```

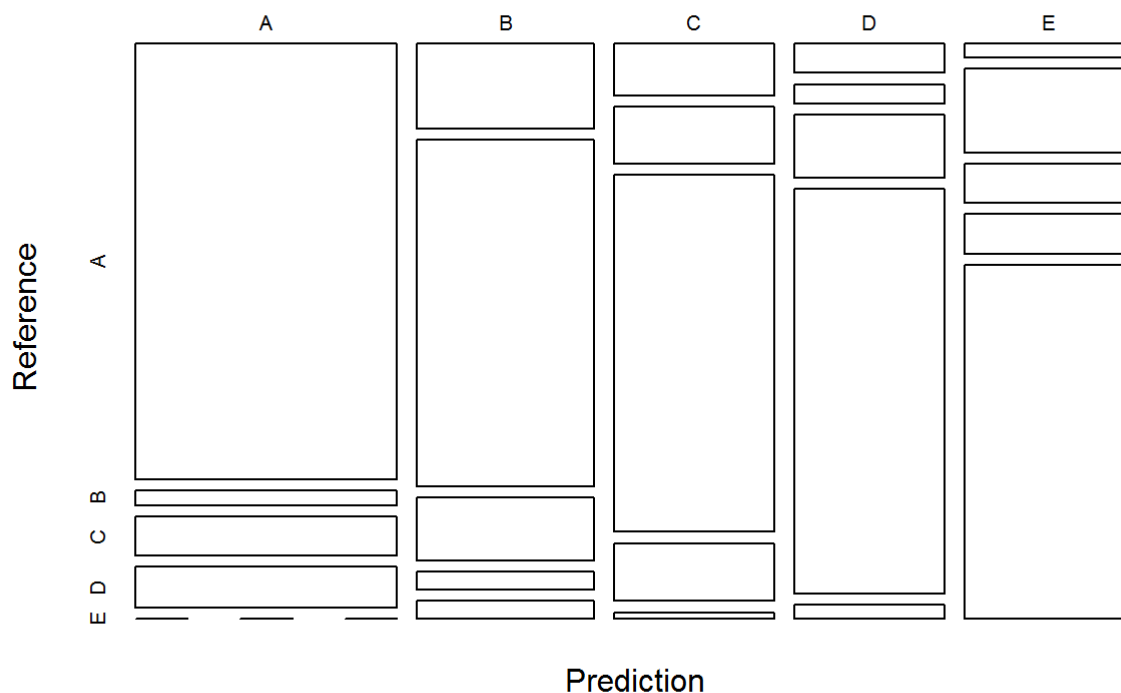
```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1377   46  123  128    0
##           B  182  745  135   38   39
##           C  100  110  692  111   13
##           D   53   35  114  736   26
##           E   28  171   79   81  723
##
## Overall Statistics
##
##           Accuracy : 0.7261
##           95% CI : (0.7145, 0.7374)
##           No Information Rate : 0.2957
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6533
##           McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.7914   0.6730   0.6054   0.6728   0.9026
## Specificity           0.9283   0.9175   0.9296   0.9524   0.9294
## Pos Pred Value        0.8226   0.6541   0.6745   0.7635   0.6682
## Neg Pred Value        0.9138   0.9237   0.9072   0.9273   0.9838
## Prevalence            0.2957   0.1881   0.1942   0.1859   0.1361
## Detection Rate        0.2340   0.1266   0.1176   0.1251   0.1229
## Detection Prevalence  0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy      0.8599   0.7953   0.7675   0.8126   0.9160
```

```
(accuracy_lda <- conf_lda$overall[1])
```

```
## Accuracy
## 0.7260833
```

```
plot(conf_lda$table, col = conf_lda$byClass, main = paste("LDA Confusion Matrix: Accuracy =",
  round(conf_lda$overall['Accuracy'], 4)))
```

LDA Confusion Matrix: Accuracy = 0.7261



Classification Tree

Next, we try the Classification Tree method (rpart)

```
modFit_rpart <- train(classe ~ ., data = training1, method = "rpart",
                      trControl = control)
print(modFit_rpart, digits = 4)
```

```
## CART
##
## 13737 samples
##   53 predictor
##   5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10991, 10991, 10989, 10988, 10989
## Resampling results across tuning parameters:
##
##   cp      Accuracy  Kappa
##   0.02777  0.5754    0.4576
##   0.04315  0.4959    0.3406
##   0.11474  0.3316    0.0721
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.02777.
```

```
predict_rpart <- predict(modFit_rpart, testing1)
(conf_rpart <- confusionMatrix(testing1$classe, predict_rpart))
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1044  179  318  130   3
##           B  206  633  245   55   0
##           C   37   44  812  133   0
##           D   45  151  506  262   0
##           E   14  264  259   53  492
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.5511
##           95% CI : (0.5382, 0.5638)
##           No Information Rate : 0.3636
##           P-Value [Acc > NIR] : < 2.2e-16
```

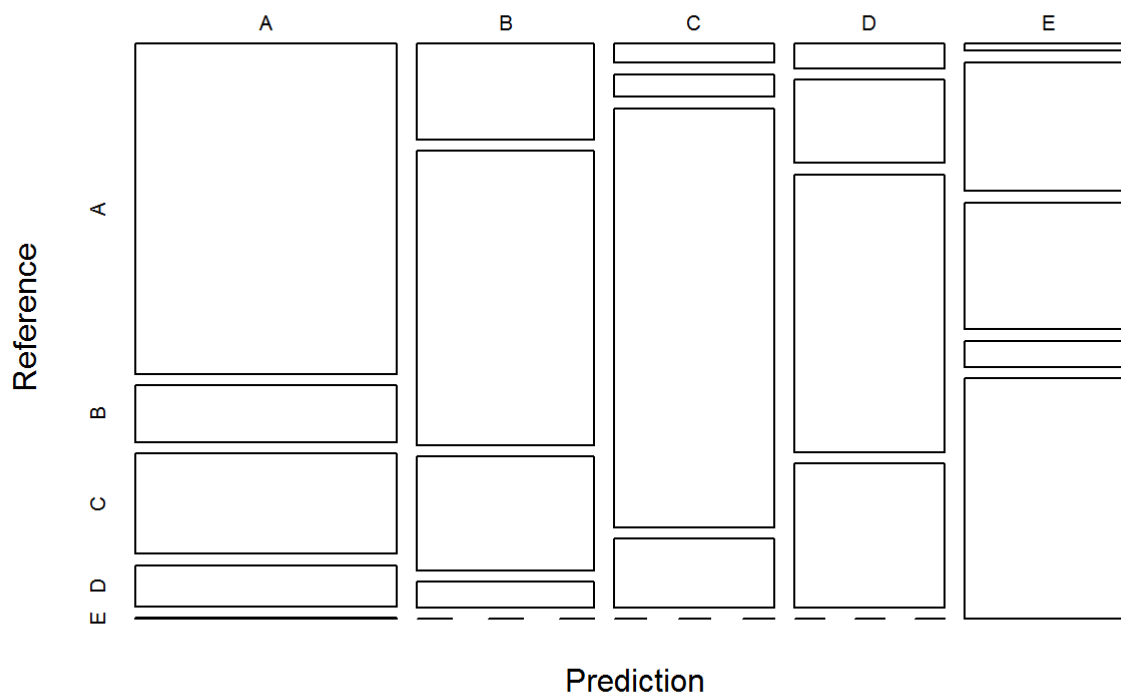
```
##
##           Kappa : 0.4365
##           McNemar's Test P-Value : < 2.2e-16
```

```
##
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity       0.7756  0.4980  0.3794  0.41390  0.99394
## Specificity       0.8612  0.8903  0.9429  0.86634  0.89054
## Pos Pred Value    0.6237  0.5558  0.7914  0.27178  0.45471
## Neg Pred Value    0.9283  0.8656  0.7267  0.92461  0.99938
## Prevalence        0.2287  0.2160  0.3636  0.10756  0.08411
## Detection Rate    0.1774  0.1076  0.1380  0.04452  0.08360
## Detection Prevalence 0.2845  0.1935  0.1743  0.16381  0.18386
## Balanced Accuracy  0.8184  0.6942  0.6611  0.64012  0.94224
```

```
plot(conf_rpart$table, col = conf_rpart$byClass, main = paste("Classification Tree Confusion
Matrix: Accuracy =", round(conf_rpart$overall['Accuracy'], 4)))
```


Classification Tree Confusion Matrix: Accuracy = 0.5511



Random Forest

Lastly, we try random forest

```
modFit_rf <- train(classe ~., data = training1, method = "rf", trControl=control )
print(modFit_rf, digits = 4)
```

```
## Random Forest
##
## 13737 samples
##    53 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10990, 10990, 10990, 10988, 10990
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.9895   0.9867
##   29    0.9908   0.9884
##   57    0.9877   0.9844
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 29.
```

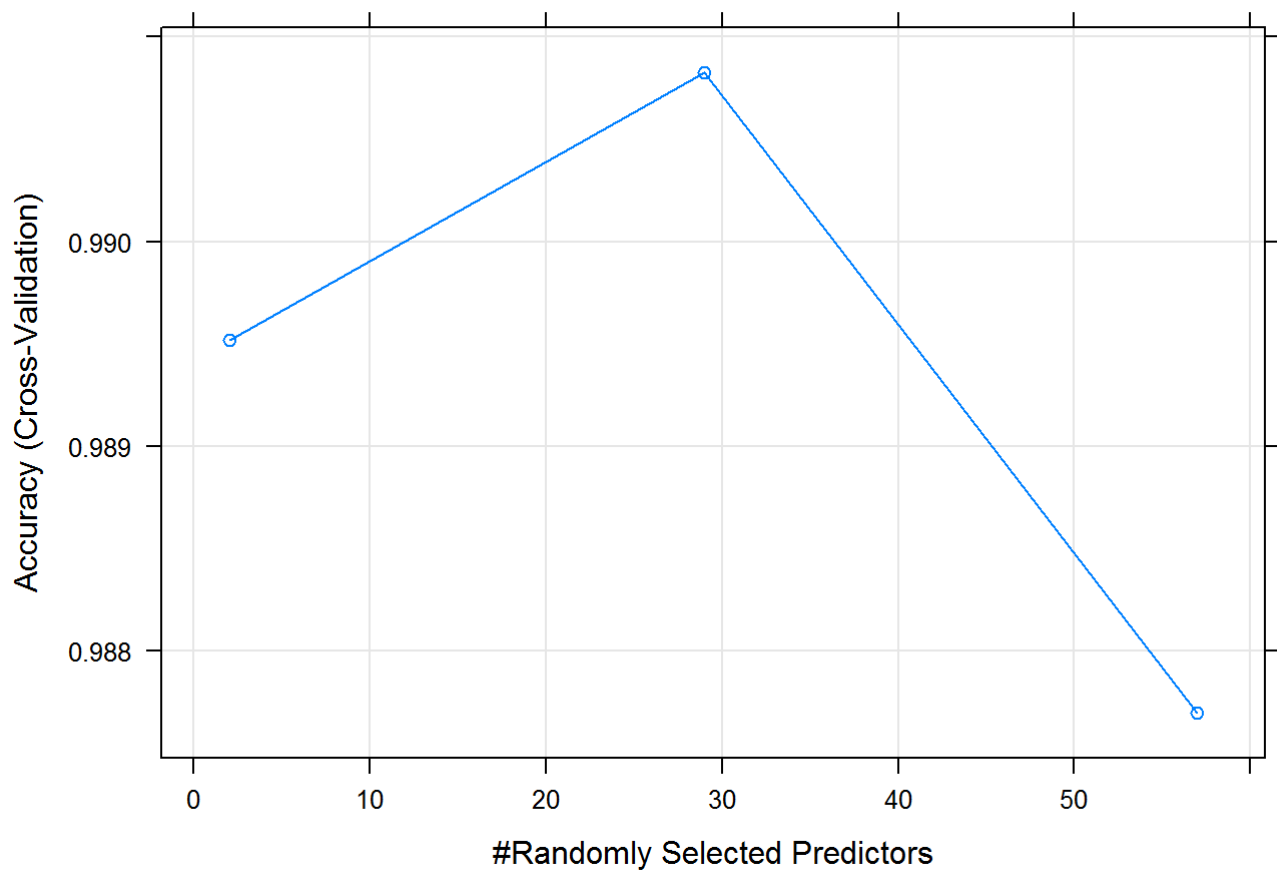
```
# predict outcomes using validation set
predict_rf <- predict(modFit_rf, testing1)
# Show prediction result
(conf_rf <- confusionMatrix(testing1$classe, predict_rf))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1674    0    0    0    0
##           B   10 1127    2    0    0
##           C    0    6 1016    4    0
##           D    0    0   11  953    0
##           E    0    0    0    0 1082
##
## Overall Statistics
##
##           Accuracy : 0.9944
##           95% CI : (0.9921, 0.9961)
##           No Information Rate : 0.2862
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9929
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9941  0.9947  0.9874  0.9958  1.0000
## Specificity      1.0000  0.9975  0.9979  0.9978  1.0000
## Pos Pred Value   1.0000  0.9895  0.9903  0.9886  1.0000
## Neg Pred Value   0.9976  0.9987  0.9973  0.9992  1.0000
## Prevalence       0.2862  0.1925  0.1749  0.1626  0.1839
## Detection Rate   0.2845  0.1915  0.1726  0.1619  0.1839
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9970  0.9961  0.9927  0.9968  1.0000
```

```
(accuracy_rf <- conf_rf$overall[1])
```

```
## Accuracy
## 0.9943925
```

```
plot(modFit_rf)
```



So, from the three models (LDA, Classification Tree, Random Forest) The accuracies are as follow

LDA: 72.6%

Classification Tree: 55%

Random Forest: 99%

As you can see, random forest so far has the best accuracy. The prediction of classe on testing dataset as follow

```
(predict(modFit_rf, testing))
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```